

```
In [64]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings('ignore')
```

```
In [65]: df = pd.read_csv("C:/Users/KahindiE/Documents/student-mat - student-mat.csv")
df.head()
```

```
Out[65]:
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	famrel	freeti
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	4	
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	5	
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	4	
3	GP	F	15	U	GT3	T	4	2	health	services	...	3	
4	GP	F	16	U	GT3	T	3	3	other	other	...	4	

5 rows × 33 columns

```
In [66]: df.columns
```

```
Out[66]: Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
              'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
              'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
              'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
              'Walc', 'health', 'absences', 'G1', 'G2', 'G3'],
              dtype='object')
```

In [67]: `df.info() # 395 rows and 33 columns`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
#   Column          Non-Null Count  Dtype
---  -
0   school          395 non-null   object
1   sex             395 non-null   object
2   age            395 non-null   int64
3   address         395 non-null   object
4   famsize         395 non-null   object
5   Pstatus         395 non-null   object
6   Medu            395 non-null   int64
7   Fedu            395 non-null   int64
8   Mjob            395 non-null   object
9   Fjob            395 non-null   object
10  reason          395 non-null   object
11  guardian        395 non-null   object
12  traveltime      395 non-null   int64
13  studytime       395 non-null   int64
14  failures        395 non-null   int64
15  schoolsup       395 non-null   object
16  famsup          395 non-null   object
17  paid            395 non-null   object
18  activities      395 non-null   object
19  nursery         395 non-null   object
20  higher          395 non-null   object
21  internet        395 non-null   object
22  romantic        395 non-null   object
23  famrel          395 non-null   int64
24  freetime        395 non-null   int64
25  goout           395 non-null   int64
26  Dalc            395 non-null   int64
27  Walc            395 non-null   int64
28  health          395 non-null   int64
29  absences        395 non-null   int64
30  G1              395 non-null   int64
31  G2              395 non-null   int64
32  G3              395 non-null   int64
dtypes: int64(16), object(17)
memory usage: 102.0+ KB
```

In [68]: `df.shape`

Out[68]: (395, 33)

In [69]: `df.describe()` # 395 rows and 33 columns

Out[69]:

	age	Medu	Fedu	traveltime	studytime	failures	famrel	fr
<b>count</b>	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.0
<b>mean</b>	16.696203	2.749367	2.521519	1.448101	2.035443	0.334177	3.944304	3.2
<b>std</b>	1.276043	1.094735	1.088201	0.697505	0.839240	0.743651	0.896659	0.9
<b>min</b>	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000	1.0
<b>25%</b>	16.000000	2.000000	2.000000	1.000000	1.000000	0.000000	4.000000	3.0
<b>50%</b>	17.000000	3.000000	2.000000	1.000000	2.000000	0.000000	4.000000	3.0
<b>75%</b>	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	5.000000	4.0
<b>max</b>	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	5.000000	5.0

In [70]: `df.isnull().sum()` # 0 null values

Out[70]:

school	0
sex	0
age	0
address	0
famsize	0
Pstatus	0
Medu	0
Fedu	0
Mjob	0
Fjob	0
reason	0
guardian	0
traveltime	0
studytime	0
failures	0
schoolsup	0
famsup	0
paid	0
activities	0
nursery	0
higher	0
internet	0
romantic	0
famrel	0
freetime	0
goout	0
Dalc	0
Walc	0
health	0
absences	0
G1	0
G2	0
G3	0
dtype:	int64

```
In [71]: #duplicates
df.duplicated().sum()
```

Out[71]: 0

```
In [72]: df.duplicated()
```

Out[72]:

0	False
1	False
2	False
3	False
4	False
...	
390	False
391	False
392	False
393	False
394	False

Length: 395, dtype: bool

```
In [73]: #1.select column Medu - mother's education (numeric: 0 - none, 1 - primary education)
df['Medu'] # 0, 1, 2, 3, 4
```

Out[73]:

0	4
1	1
2	1
3	4
4	3
..	
390	2
391	3
392	1
393	3
394	1

Name: Medu, Length: 395, dtype: int64

```
In [74]: #2. Fedu - father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - high school, 3 - college)
df['Fedu'] # 0, 1, 2, 3, 4
```

Out[74]:

0	4
1	1
2	1
3	2
4	3
..	
390	2
391	1
392	1
393	2
394	1

Name: Fedu, Length: 395, dtype: int64

```
In [75]: #3. famrel - quality of family relationships (numeric: from 1 - very bad to 5 -  
df['famrel'] # 1, 2, 3, 4, 5
```

```
Out[75]: 0      4  
1      5  
2      4  
3      3  
4      4  
      ..  
390    5  
391    2  
392    5  
393    4  
394    3  
Name: famrel, Length: 395, dtype: int64
```

```
In [76]: #4. freetime - free time after school (numeric: from 1 - very low to 5 - very h  
df['freetime']
```

```
Out[76]: 0      3  
1      3  
2      3  
3      2  
4      3  
      ..  
390    5  
391    4  
392    5  
393    4  
394    2  
Name: freetime, Length: 395, dtype: int64
```

```
In [77]: #5. goout - going out with friends (numeric: from 1 - very low to 5 - very high  
df['goout']
```

```
Out[77]: 0      4  
1      3  
2      2  
3      2  
4      2  
      ..  
390    4  
391    5  
392    3  
393    1  
394    3  
Name: goout, Length: 395, dtype: int64
```

```
In [78]: #6. Dalc - workday alcohol consumption (numeric: from 1 - very low to 5 - very  
df['Dalc'] # 1, 2, 3, 4, 5
```

```
Out[78]: 0      1  
        1      1  
        2      2  
        3      1  
        4      1  
        ..  
       390     4  
       391     3  
       392     3  
       393     3  
       394     3  
Name: Dalc, Length: 395, dtype: int64
```

```
In [79]: #7. Walc - weekend alcohol consumption (numeric: from 1 - very low to 5 - very  
df['Walc']
```

```
Out[79]: 0      1  
        1      1  
        2      3  
        3      1  
        4      2  
        ..  
       390     5  
       391     4  
       392     3  
       393     4  
       394     3  
Name: Walc, Length: 395, dtype: int64
```

```
In [80]: #8. health - current health status (numeric: from 1 - very bad to 5 - very good  
df['health'] # 1, 2, 3, 4, 5
```

```
Out[80]: 0      3  
        1      3  
        2      3  
        3      5  
        4      5  
        ..  
       390     4  
       391     2  
       392     3  
       393     5  
       394     5  
Name: health, Length: 395, dtype: int64
```

```
In [81]: #9. absences - number of school absences (numeric: from 0 to 93)
df['absences']
```

```
Out[81]: 0      6
         1      4
         2     10
         3      2
         4      4
         ..
        390    11
        391      3
        392      3
        393      0
        394      5
        Name: absences, Length: 395, dtype: int64
```

```
In [82]: #10. G3 - final grade (numeric: from 0 to 20, output target)
df['G3'] # 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19
```

```
Out[82]: 0      6
         1      6
         2     10
         3     15
         4     10
         ..
        390      9
        391     16
        392      7
        393     10
        394      9
        Name: G3, Length: 395, dtype: int64
```

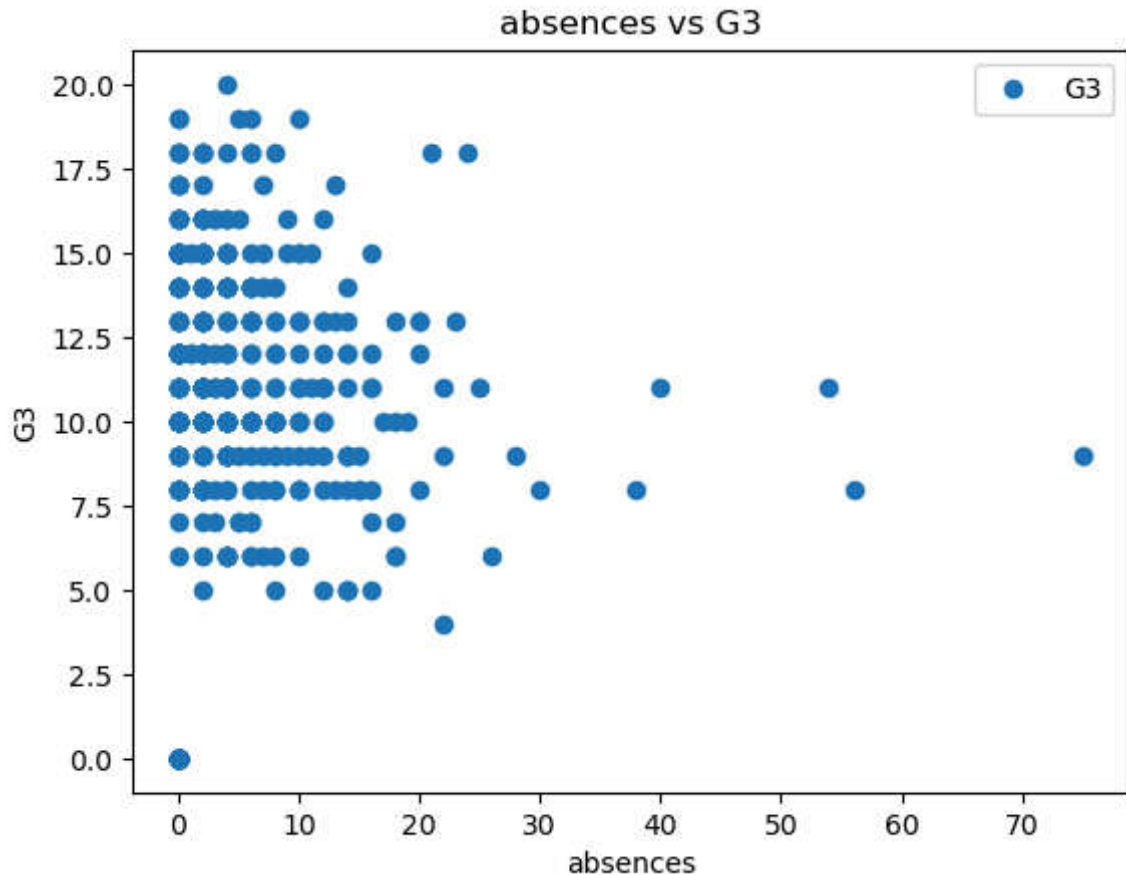
```
In [83]: #handling missing values  
df.isnull().sum() # 0 null values
```

```
Out[83]: school      0  
sex                0  
age               0  
address           0  
famsize           0  
Pstatus           0  
Medu              0  
Fedu              0  
Mjob              0  
Fjob              0  
reason            0  
guardian          0  
traveltime        0  
studytime         0  
failures          0  
schoolsup         0  
famsup            0  
paid              0  
activities        0  
nursery           0  
higher            0  
internet          0  
romantic          0  
famrel            0  
freetime          0  
goout             0  
Dalc              0  
Walc              0  
health            0  
absences          0  
G1                0  
G2                0  
G3                0  
dtype: int64
```



```
In [101]: #scatter plot
df.plot(x='absences', y='G3', style='o')
plt.title('absences vs G3')
plt.xlabel('absences')
plt.ylabel('G3')
```

Out[101]: Text(0, 0.5, 'G3')



```
In [96]: # Carry out multiple linear regression analysis to predict G3 - final grade (ou
# 1. Select the independent variables (X) and the dependent variable (y)
import statsmodels.api as sm
X = df[['Medu', 'Fedu', 'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health']
y = df['G3'] # dependent variable
X = sm.add_constant(X) # adding a constant
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random
```

```
In [85]: # 2. Create an instance of the LinearRegression model
model = sm.OLS(y, X).fit() # Ordinary Least Squares
model.summary() # Summary of the model
```

Out[85]: OLS Regression Results

<b>Dep. Variable:</b>	G3	<b>R-squared:</b>	0.083
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.062
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	3.870
<b>Date:</b>	Tue, 11 Jun 2024	<b>Prob (F-statistic):</b>	0.000103
<b>Time:</b>	15:54:50	<b>Log-Likelihood:</b>	-1144.1
<b>No. Observations:</b>	395	<b>AIC:</b>	2308.
<b>Df Residuals:</b>	385	<b>BIC:</b>	2348.
<b>Df Model:</b>	9		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	8.4165	1.474	5.711	0.000	5.519	11.314

```
In [86]: # Evaluating the model
y_pred = model.predict(X_test) # Predicted values
print('Mean Squared Error:', mean_squared_error(y_test, y_pred))
print('Mean Absolute Error:', mean_absolute_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(mean_squared_error(y_test, y_pred)))
print('R-squared:', r2_score(y_test, y_pred))
print('Adjusted R-squared:', 1 - (1-r2_score(y_test, y_pred))*(len(y)-1)/(len(y)-10))
```

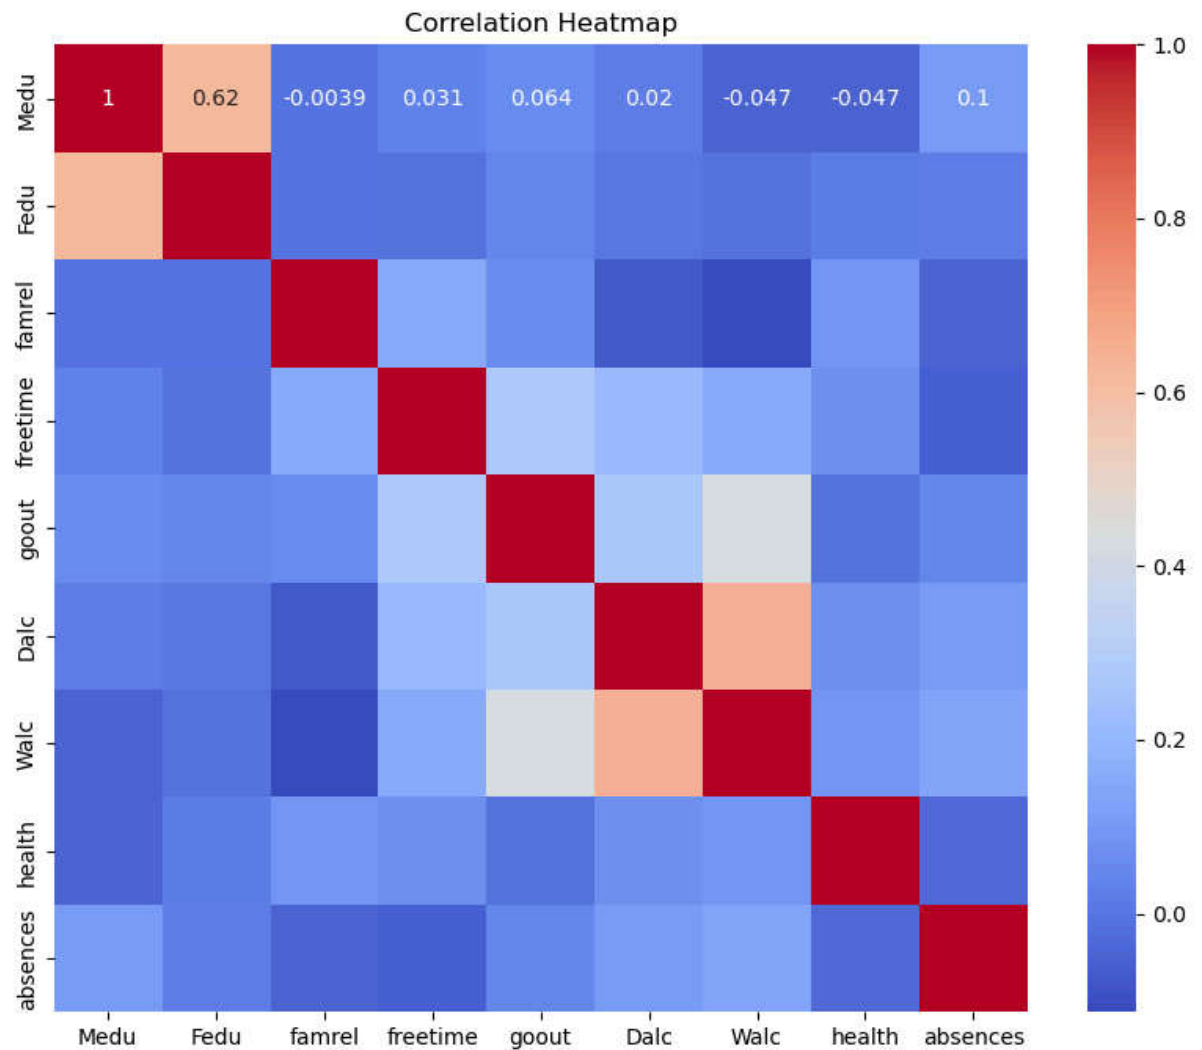
Mean Squared Error: 19.402058229506007  
Mean Absolute Error: 3.4955840535276566  
Root Mean Squared Error: 4.404776751380938  
R-squared: 0.11736865350403947  
Adjusted R-squared: 0.09438346218904048

```
In [87]: #define cleaned_df dataframe
cleaned_df = df[['Medu', 'Fedu', 'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences'], cleaned_df.head()
```


Out[87]:

	Medu	Fedu	famrel	freetime	goout	Dalc	Walc	health	absences
0	4	4	4	3	4	1	1	3	6
1	1	1	5	3	3	1	1	3	4
2	1	1	4	3	2	2	3	3	10
3	4	2	3	2	2	1	1	5	2
4	3	3	4	3	2	1	2	5	4

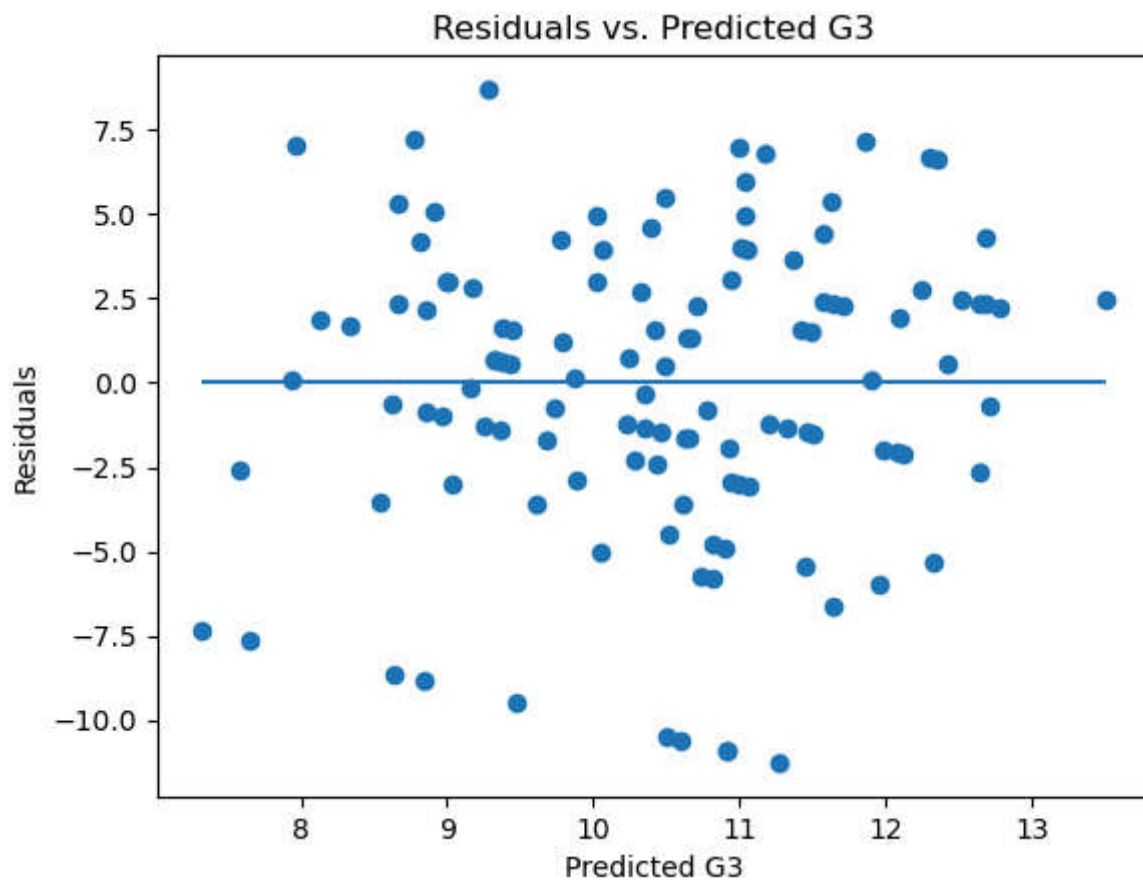
```
In [88]: #vizualization with heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(cleaned_df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```




```
In [89]: #Test whether all the multiple linear regression assumptions are met.  
plt.scatter(y_test, y_pred)  
plt.xlabel('Actual G3')  
plt.ylabel('Predicted G3')  
plt.title('Actual vs. Predicted G3')  
plt.show()
```



```
In [90]: #2. Homoscedasticity
plt.scatter(y_pred, y_test - y_pred)
plt.hlines(y=0, xmin=y_pred.min(), xmax=y_pred.max())
plt.xlabel('Predicted G3')
plt.ylabel('Residuals')
plt.title('Residuals vs. Predicted G3')
plt.show()
```



```
In [91]: #3. Normality
import scipy.stats as stats # Importing the stats module from the scipy library
residuals = y_test - y_pred # Calculating the residuals
stats.probplot(residuals, dist="norm", plot=plt) # Plotting the normal probability plot
plt.title("Normal Q-Q plot")
plt.show()
```



```
In [92]: #4. Multicollinearity
vif = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])] # Calculating VIF values
pd.DataFrame({'VIF': vif}, index=X.columns) # VIF values for each independent variable
```

Out[92]:

	VIF
<b>const</b>	43.554418
<b>Medu</b>	1.691386
<b>Fedu</b>	1.651711
<b>famrel</b>	1.063832
<b>freetime</b>	1.156675
<b>goout</b>	1.325336
<b>Dalc</b>	1.778884
<b>Walc</b>	2.033048
<b>health</b>	1.034931
<b>absences</b>	1.044119

```
In [93]: #Thanks finally I reached the end of the Assigment.
```